

Dynamic Network Traffic Monitoring and Scaling Using eBPF

Team Members:

Faiaz Halim (202383578)

Anil Manyam (202483873)

S. M. Yeamin Oni (202293272)

Problem Statement

Traditional monitoring lacks flexibility for microservices

Need for real-time traffic-based scaling

Complex network policy management

Solution

eBPF-based network monitoring

Multi-metric scaling (traffic, CPU, memory)

Policy-driven traffic control

Real-time visualization

Architecture

Components

1. Core Infrastructure

- Kubernetes 1.31 cluster (kind)
- Cilium 1.16.3 with eBPF
- Hubble for flow visibility

2. Monitoring Stack

- Prometheus for metrics storage
- Custom Prometheus Adapter
- Hubble UI and Grafana for visualization

3. Application Layer

- Frontend service
- Backend service
- Load testing suite

Implementation

1. Network Monitoring

- Cilium eBPF programs for packet inspection
- Hubble flow monitoring

```
apiVersion: cilium.io/v2
kind: CiliumClusterwideNetworkPolicy
metadata:
  name: backend-access
spec:
  endpointSelector:
    matchLabels:
      app: backend
  ingress:
    - fromEndpoints:
      - matchLabels:
          app: frontend
```

2. Dynamic Scaling

- Multiple HPA metrics combined:

```
spec:
  metrics:
    - type: Object
      object:
        metric:
          name: http_requests_per_second
    - type: Resource
      resource:
        name: cpu
    - type: Resource
      resource:
        name: memory
```

3. Load Testing Suite

- Traffic generation
- CPU stress testing
- Real-time monitoring
- Automated reporting

Features

Monitoring

- Network flows
- Resource utilization
- Policy enforcement

Scaling

- Traffic-based scaling
- Resource-based scaling
- Combined metric approach

Security and Control

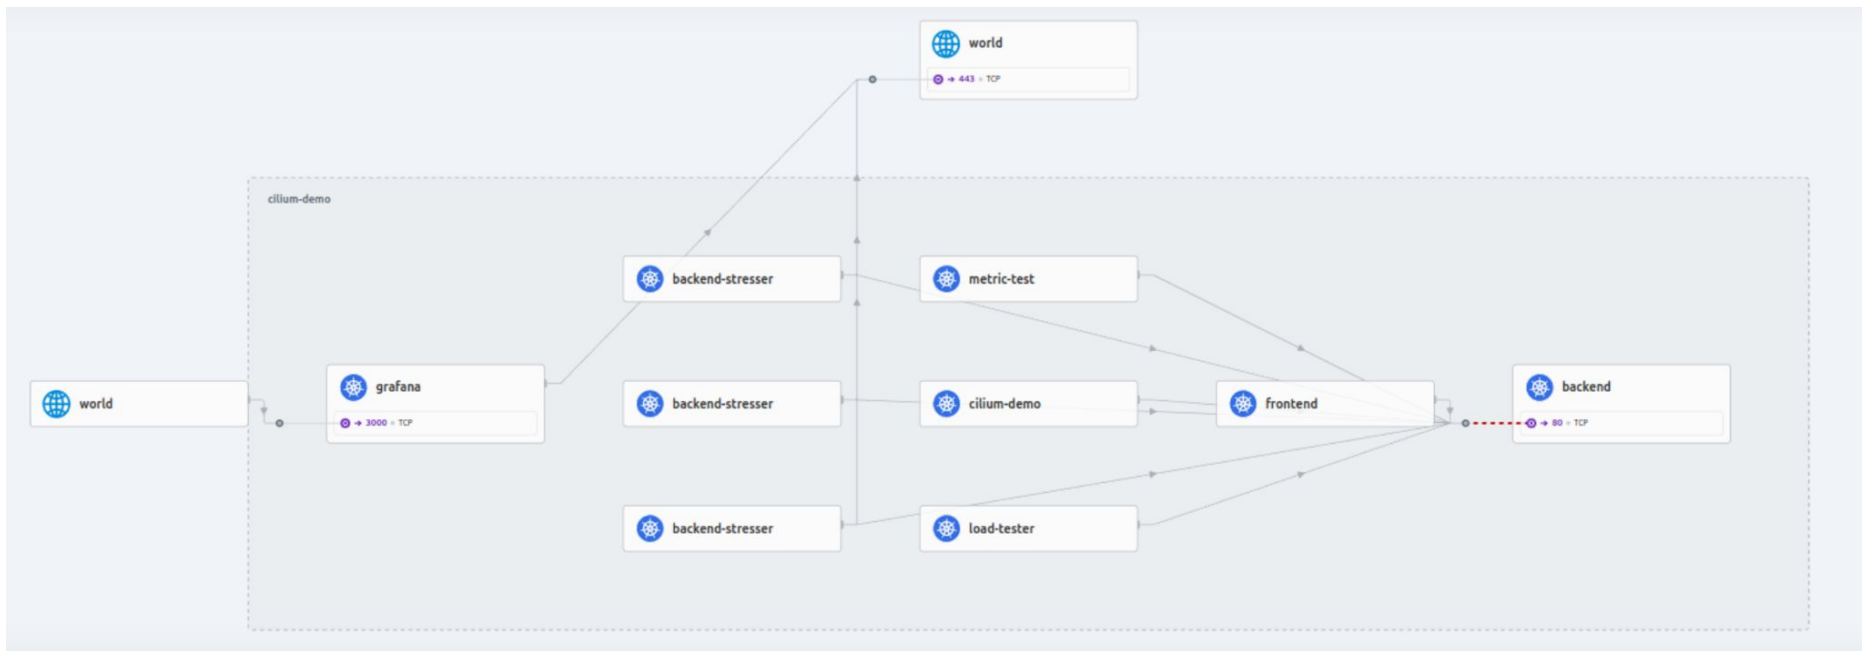
- L3/L4/L7 policies
- Rate limiting
- Flow visibility

Testing and Validation

- Automated load testing
- Multi-metric validation
- Comprehensive reporting

Results

- Hubble UI
- Traffic Monitoring Report



Future Enhancements

Advanced Analytics

- Machine learning for predictive scaling
- Pattern recognition
- Anomaly detection

Enhanced Policies

- Dynamic policy adjustment
- Context-aware rules
- Advanced rate limiting

Optimization

- Performance tuning
- Resource efficiency
- Scaling algorithms

Questions?

References

- S. Magnani, F. Risso, D. Siracusa, “*A Control Plane Enabling Automated and Fully Adaptive Network Traffic Monitoring With eBPF*”, IEEE Access (Volume: 10) (<https://ieeexplore.ieee.org/document/9869628>), 2022.
- JB. Lee, TH. Yoo, EH. Lee, BH. Hwang, SW. Ahn, CH. Cho, “*High-Performance Software Load Balancer for Cloud-Native Architecture*”, IEEE Access (Volume: 9) (<https://ieeexplore.ieee.org/document/9524915>), 2021.
- S. Miano, F. Risso, M. V. Bernal, M. Bertrone, Y. Lu, “*A Framework for eBPF-Based Network Functions in an Era of Microservices*”, IEEE Transactions on Network and Service Management (Volume: 18, Issue: 1) (<https://ieeexplore.ieee.org/document/9340283>), 2021.
- H. Sharaf, I. Ahmad, T. Dimitriou, “*Extended Berkeley Packet Filter: An Application Perspective*”, IEEE Access (Volume: 10) (<https://ieeexplore.ieee.org/document/9968265>), 2022.
- M. Jadin, Q. D. Coninck, L. Navarre, M. Schapira, O. Bonaventure, “*Leveraging eBPF to Make TCP Path-Aware*”, IEEE Transactions on Network and Service Management (Volume: 19, Issue: 3) (<https://ieeexplore.ieee.org/document/9772044>), 2022.